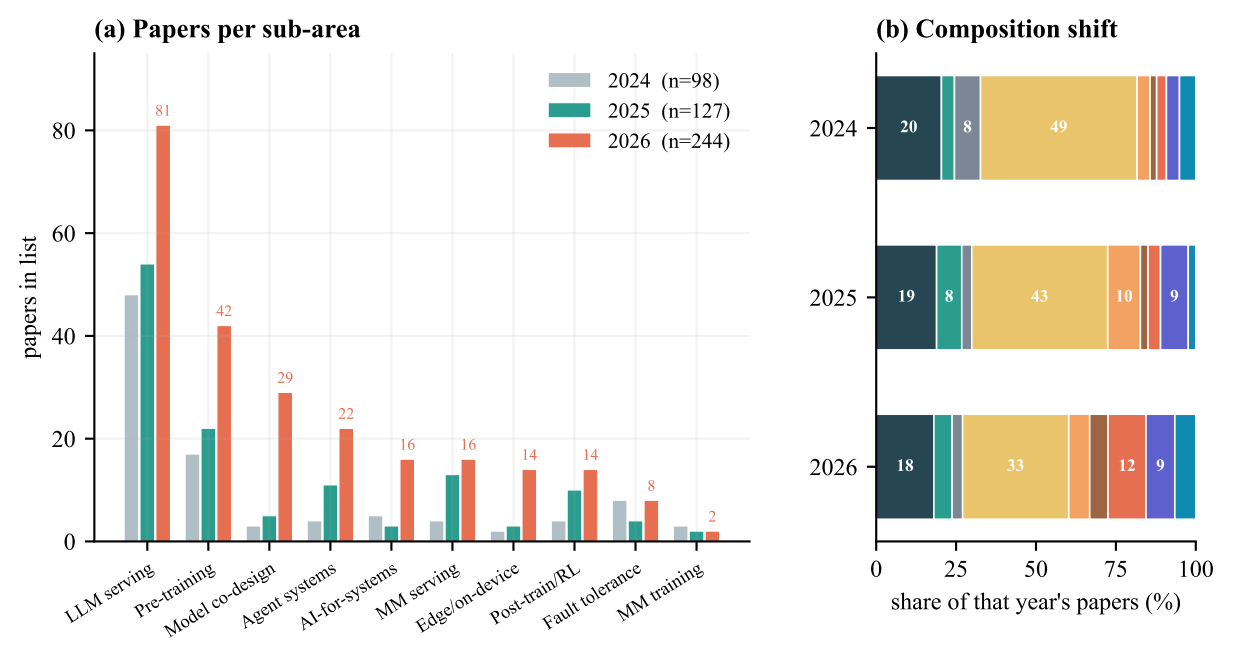
(a) Papers per sub-area
2024 (n=98)
2025 (n=127)
2026 (n=244)
papers in list
0
20
40
60
80
81
42
29
22
16
16
14
14
8
2
LLM serving
Pre-training
Model co-design
Agent systems
AI-for-systems
MM serving
Edge/on-device
Post-train/RL
Fault tolerance
MM training
(b) Composition shift
2024
2025
2026
20
8
49
19
8
43
10
9
18
33
12
9
0
25
50
75
100
share of that year's papers (%)
Training infra
Post-training / RL
Reliability
Serving runtime
Multi-modal / diffusion
Edge / on-device
Kernel / model co-design
Agentic systems
AI-for-systems